Hi team,

I build the infrastructure that takes GenAI from demo to production at scale. Adobe's push into agentic AI for enterprise content orchestration, and the Data Platforms team's role in powering that, is exactly the kind of problem I want to solve next.

At Bigstep Technologies, I lead GenAI platform delivery as Lead Architect across clients ranging from startups to global enterprises. The work that maps most directly to this role: I designed a **Field-Aware Extraction Engine processing 2M+ records/month**, routing each field by complexity so ~60% bypass the LLM entirely, cutting token spend ~40% without accuracy loss. I also built a **RAG pipeline fusing SPLADE and Sentence-BERT retrieval with Cross-Encoder re-ranking**, lifting NDCG@5 from 0.71 to 0.86. Both systems run on Python, Docker, and AWS, with CI/CD evaluation gates that block deploys when quality metrics regress.

Before that, as Founding Engineer at Kuttl, I **fine-tuned Llama-3-8B (LoRA, Hugging Face Transformers)** to replace GPT-4 for structured extraction, achieving 98% schema compliance while cutting inference costs by 90%. I built the distributed data workflows (FastAPI, Temporal.io, 50k+ messages/day) and operationalized model evaluation with RAGAS scoring in CI/CD, catching regressions before production. At Amazon, I owned high-scale microservices serving 3,000+ merchants at 99.99% availability and reduced cloud infrastructure costs by **$100k/year** through Spot Instance re-architecture.

What draws me to Adobe specifically is the scale and ambition: shipping AI agents that reshape how businesses orchestrate customer experiences requires the intersection of scalable data engineering, reliable agentic systems, and cost-conscious LLMOps. That intersection is where I operate best.

I would welcome the chance to discuss how my experience translates to the Data Platforms team's goals.

Best, Parvez Akhtar parvezakhtar218@gmail.com | +91 888-232-7732