Hi team,

I have spent the last two years building the unglamorous parts of AI agents that the demo never shows: keeping context coherent across large inputs, choosing the right model per task, and catching regressions before users do. Turning natural language into production-ready full-stack apps is exactly that problem at the hardest setting, and it is the one I want to work on at Bolt.new.

At VTS, I built an agentic document-intelligence pipeline that abstracts structured terms from commercial real-estate leases with structured-output validation. When field-level accuracy dropped 10 points, I treated it as a data problem: I analyzed failure modes across multi-turn extraction, then redesigned the prompt strategy and image preprocessing to recover the full drop and restore the roughly 88% baseline. That loop of measure, diagnose, and iterate is how I think agents get reliable.

Context and cost are the other half. On a GenAI platform processing 2M+ records a month, I built a Strategy Router that classifies each field by complexity and routes about 60% to deterministic parsers, reserving the LLM for the rest and cutting token spend roughly 40% with no measurable accuracy loss. I have orchestrated multi-step, multi-agent workflows with tool and function calling over MCP, and I built CI evaluation gates (RAGAS) that block deploys when faithfulness or answer-relevancy slips. I have also chosen models on evidence, fine-tuning Llama-3-8B to replace GPT-4 for a structured task after benchmarking, holding 98% schema compliance while cutting per-token cost 90%.

What pulls me to Bolt is the pace and the stakes: an agent shipping to millions, where context management and tool orchestration are the product, not a feature. I work daily across Python and TypeScript and like owning problems end to end. I would love to help push what these agents can reliably build.

Best, Parvez Akhtar +91 888-232-7732 | parvezakhtar218@gmail.com