

Hi team,

I build AI automation that survives production, not just demos, which is exactly what scaling a mission-critical subscriptions platform demands. As Lead Architect at Bigstep Technologies, I design the full path from ingestion through orchestration, model inference, human-in-the-loop review, and monitoring, so the Senior AI Engineer role on your intelligent automation team maps closely to what I already do.

Most relevant to this role: I orchestrated a **serverless multi-agent framework for Salesforce (CRM) automation** using **tool/function calling over MCP (Model Context Protocol)**, wiring LLM inference into enterprise workflows and cutting idle compute with a pure pay-per-execution model that **saved roughly \$45k/year**. On a GenAI platform for Outcomelabs.ai, I architected a field-aware extraction engine processing **2M+ records/month**, with a strategy router that sends about 60% of fields to deterministic parsers and reserves the LLM for the rest, **trimming token spend ~40% with no accuracy loss**. That blend of LLMs, classical ML, and RPA-style workflow automation is the core of your job description.

The subscriptions focus resonates with me. At Amazon I co-developed the SmartPOS **billing engine, scaling it to 3,000+ enterprise merchants at 99.99% availability**, so I understand what billing and lifecycle systems demand under real load. I also ship the guardrails these systems need: structured-output validation, prompt-injection defense, PII-aware handling, and **CI/CD evaluation gates (RAGAS)** that block regressions before release.

GE HealthCare's goal of a world where healthcare has no limits is the rare place where automation directly frees clinicians for patient care. I would welcome the chance to help shape that platform from the ground up, across **AWS Bedrock and Azure OpenAI**, with the reliability and cost discipline it deserves.

Best, Parvez Akhtar +91 888-232-7732 | parvezakhtar218@gmail.com