Hi team,

The hardest part of enterprise AI is not the model, it is shipping it reliably, securely, and cheaply enough to trust in production. That is the work I do, and doing it inside a regulated bank scaling agentic and generative AI through a dedicated AI Centre of Excellence is exactly the challenge I want. Lloyds Technology Centre's focus on operationalising AI safely maps directly to how I already build.

As Lead Architect at Bigstep, I run AI as an engineering discipline. I built **CI/CD evaluation gates** with **Ragas** graders scoring faithfulness and answer-relevancy that automatically block deployments below threshold, so quality regressions never reach production. On a **2M+ records/month** extraction engine, I designed a **Strategy Router** that tracks latency, token usage, and cost and routes ~60% of fields to deterministic parsers, cutting **token spend ~40%** with no accuracy loss. Earlier, I fine-tuned **Llama-3-8B (LoRA/PEFT)** to hit **98% schema compliance** while reducing per-token cost **90%**.

On the engineering and security side, I ship models as **FastAPI microservices** and have built enterprise access controls end to end: a **Keycloak** and **Spring Security** identity broker using stateless **JWT**, and a **Schema-per-Tenant RBAC** architecture isolating 21 enterprise clients with automated data-integrity reconciliation. At Amazon I held **99.99% availability** on billing microservices and cut infrastructure cost **25%**. This is the reliability, security, and cost discipline your role centres on.

What draws me to Lloyds specifically is the commitment to Responsible AI: building GenAI that 27 million customers can trust is a governance problem as much as a modelling one, and I want to set those technical standards and mentor engineers doing it. I would welcome the chance to discuss how I can contribute to the AI CoE.

Best, Parvez Akhtar +91 888-232-7732 | parvezakhtar218@gmail.com