

Hi team,

I build backend systems that take GenAI from prototype to dependable production, which is exactly what PwC's Acceleration Centers ask of a Senior Associate on the Data and Analytics team. As Lead Architect at Bigstep, I already do this for global enterprises including **CNBC London** and **VTs**, translating ambiguous business goals into **event-driven, agentic** architectures that clients run in production.

Your role centers on scalable backend solutions, agentic AI, and clean deployment, and my recent work maps directly to it. I orchestrated a **serverless, event-driven multi-agent** framework for Salesforce automation using **tool/function calling over MCP (Model Context Protocol)**, where a pay-per-execution model **cut idle compute and saved ~\$45k/year**. On a GenAI extraction platform, I architected a **Field-Aware Extraction Engine processing 2M+ records/month** and a routing layer that **cut token spend ~40%** with no measurable accuracy loss. Earlier, I built a distributed document pipeline on **FastAPI and Temporal.io** handling **50k+ messages/day** with retry logic and dead-letter queues so no data is lost.

The stack you describe is the one I work in daily: **Python (FastAPI), LangChain and LlamaIndex, PostgreSQL and MongoDB, Kubernetes and Docker**, and **CI/CD** across **AWS and Azure**. I also care about the engineering underneath, applying **SOLID and clean-architecture** principles when I led a Strangler-Fig migration to Spring Boot microservices that **cut cyclomatic complexity ~42%**. I partner naturally with product, ML, and DevOps teams because shipping reliable AI for clients depends on that collaboration.

What draws me to PwC specifically is the range: solving GenAI problems across many clients and industries rather than one product, with the quality bar and AC engagement demands. I would welcome the chance to bring my agentic AI and backend experience to that work.

Best, Parvez Akhtar +91 888-232-7732 | parvezakhtar218@gmail.com