Hi team,

I build the backend that turns GenAI prototypes into production systems, and PwC's Acceleration Centers are doing exactly that at enterprise scale. Your Data and Analytics team wants a Senior Associate who can translate business goals into scalable **GenAI and Agentic AI** software, working across product, ML, and DevOps. That is the work I do now, and I would like to bring it to your client engagements.

Most recently, as Lead Architect at Bigstep Technologies, I orchestrated a **serverless multi-agent** framework for Salesforce automation using tool/function calling over **MCP (Model Context Protocol)**, cutting idle compute to a pay-per-execution model that saves ~**$45k/year**. I also architected a Field-Aware Extraction Engine processing **2M+ records/month**, where a routing layer sends ~60% of fields to deterministic parsers and reserves the LLM for the rest, trimming token spend ~**40%** with no accuracy loss. These are the event-driven, API-integrated backends the role calls for.

The Python and architecture depth is genuine. At Kuttl I built a distributed document pipeline on **FastAPI and Temporal.io** handling **50k+ daily messages** with asynchronous, event-driven processing, retries, and dead-letter queues, and fine-tuned Llama-3-8B to cut inference cost **90%**. Earlier, I led a Strangler-Fig migration to Spring Boot **microservices** on **SOLID and clean-architecture** principles, and ran high-scale services on **Kubernetes** and AWS. I work fluidly across Python and Java, relational and NoSQL stores, and CI/CD.

What draws me to PwC specifically is the breadth: shipping reliable GenAI across many clients and domains, not one product, is a harder and more interesting engineering problem, and the quality bar the ACs are known for matches how I like to build.

I would welcome the chance to discuss how I can contribute.

Best, Parvez Akhtar +91 888-232-7732 | parvezakhtar218@gmail.com