

Hi team,

I build agentic AI systems that turn messy documents into clean, trustworthy structured data, and I want to do that work where it saves lives. Truveta's mission of Saving Lives with Data, and the Truveta Language Model's job of transforming multimodal EHR data into billions of clean data points, is the exact problem I have spent the last year solving in other domains. Joining the Intelligence Platform and Applications team to build adaptive, self-improving agents is the kind of purpose I am looking for.

Most recently, as Lead Architect at Bigstep, I built a **multimodal document-intelligence pipeline** for VTS that abstracts structured terms from commercial documents with structured-output validation. When accuracy dropped, I diagnosed a **10-point regression** and redesigned the prompt strategy and image preprocessing to recover the full drop back to its **~88% baseline**, the same reasoning, evaluation, and iteration loop your extraction models demand. I also architected an extraction engine processing **2M+ records/month** and orchestrated a **multi-agent** framework over **MCP (Model Context Protocol)** with tool and function calling, saving **~\$45k/year**.

On the ML side, I fine-tuned **Llama-3-8B (LoRA/PEFT)** to replace GPT-4 for structured extraction, hitting **98% schema compliance** while cutting per-token cost **90%**. I run evaluation as an engineering discipline: **CI/CD eval gates** with **LLM-as-judge** graders (Ragas) scoring faithfulness and answer-relevancy, so quality regressions never reach production. That instinct for measurable, trustworthy AI mirrors Truveta's own commitment to transparency and quality.

What draws me to Truveta specifically is that intelligence here is accountable to patients and clinicians, not just benchmarks. I would bring senior ownership, RAG and agentic depth, and a bias for shipping reliable systems to the IPA team.

I would welcome the chance to talk.

Best, Parvez Akhtar +91 888-232-7732 | parvezakhtar218@gmail.com